

```
# https://towardsdatascience.com/anomaly-detection-using-autoencoders-5b032178a1ea

import pandas as pd
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix, recall_score, accuracy_score, precision_score
RANDOM_SEED = 2021
TEST_PCT = 0.3
LABELS = ["Normal", "Fraud"]

from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive


dataset = pd.read_csv("/content/drive/MyDrive/DL assignment/DL csv folder/creditcard.csv")
```

## ▼ New Section

```
#check for any nullvalues
print("Any nulls in the dataset ",dataset.isnull().values.any() )
print('-----')
print("No. of unique labels ", len(dataset['Class'].unique()))
print("Label values ",dataset.Class.unique())
#0 is for normal credit card transaction
#1 is for fraudulent credit card transaction
print('-----')
print("Break down of the Normal and Fraud Transactions")
print(pd.value_counts(dataset['Class'], sort = True) )
```

```
Any nulls in the dataset  False
-----
No. of unique labels  2
Label values  [0 1]
-----
Break down of the Normal and Fraud Transactions
0      284315
1         492
Name: Class, dtype: int64
```

```
#Visualizing the imbalanced dataset
#plotting the number of normal and fraud transactions in the dataset.
count_classes = pd.value_counts(dataset['Class'], sort = True)
count_classes.plot(kind = 'bar', rot=0)
plt.xticks(range(len(dataset['Class'].unique()), dataset.Class.unique()))
plt.title("Frequency by observation number")
plt.xlabel("Class")
plt.ylabel("Number of Observations");
```



```
#The last column in the dataset is our target variable.
raw_data = dataset.values
# The last element contains if the transaction is normal which is represented by a 0 and 1
labels = raw_data[:, -1]
# The other data points are the electrocardriogram data
data = raw_data[:, 0:-1]
train_data, test_data, train_labels, test_labels = train_test_split(
    data, labels, test_size=0.2, random_state=2021
)
```

```
#Normalize the data to have a value between 0 and 1
min_val = tf.reduce_min(train_data)
max_val = tf.reduce_max(train_data)
train_data = (train_data - min_val) / (max_val - min_val)
test_data = (test_data - min_val) / (max_val - min_val)
train_data = tf.cast(train_data, tf.float32)
test_data = tf.cast(test_data, tf.float32)
```

```
# Use only normal transactions to train the Autoencoder.
train_labels = train_labels.astype(bool)
test_labels = test_labels.astype(bool)
#creating normal and fraud datasets
normal_train_data = train_data[~train_labels]
normal_test_data = test_data[~test_labels]
fraud_train_data = train_data[train_labels]
fraud_test_data = test_data[test_labels]
print(" No. of records in Fraud Train Data=",len(fraud_train_data))
print(" No. of records in Normal Train data=",len(normal_train_data))
```

```

print(" No. of records in Fraud Test Data=",len(fraud_test_data))
print(" No. of records in Normal Test data=",len(normal_test_data))

    No. of records in Fraud Train Data= 389
    No. of records in Normal Train data= 227456
    No. of records in Fraud Test Data= 103
    No. of records in Normal Test data= 56859

#Set the training parameter values
nb_epoch = 50
batch_size = 64
input_dim = normal_train_data.shape[1] #num of columns, 30
encoding_dim = 14
hidden_dim_1 = int(encoding_dim / 2) #
hidden_dim_2=4
learning_rate = 1e-7

#Create the Autoencoder

#input Layer
input_layer = tf.keras.layers.Input(shape=(input_dim, ))
#Encoder
encoder = tf.keras.layers.Dense(encoding_dim, activation="tanh",activity_regularizer=tf.ke
encoder=tf.keras.layers.Dropout(0.2)(encoder)
encoder = tf.keras.layers.Dense(hidden_dim_1, activation='relu')(encoder)
encoder = tf.keras.layers.Dense(hidden_dim_2, activation=tf.nn.leaky_relu)(encoder)
# Decoder
decoder = tf.keras.layers.Dense(hidden_dim_1, activation='relu')(encoder)
decoder=tf.keras.layers.Dropout(0.2)(decoder)
decoder = tf.keras.layers.Dense(encoding_dim, activation='relu')(decoder)
decoder = tf.keras.layers.Dense(input_dim, activation='tanh')(decoder)
#Autoencoder
autoencoder = tf.keras.Model(inputs=input_layer, outputs=decoder)
autoencoder.summary()

```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 30)]	0
dense (Dense)	(None, 14)	434
dropout (Dropout)	(None, 14)	0
dense_1 (Dense)	(None, 7)	105
dense_2 (Dense)	(None, 4)	32
dense_3 (Dense)	(None, 7)	35
dropout_1 (Dropout)	(None, 7)	0
dense_4 (Dense)	(None, 14)	112

dense_5 (Dense)	(None, 30)	450
-----------------	------------	-----

```
=====
Total params: 1,168
Trainable params: 1,168
Non-trainable params: 0
=====
```

---

```
# Define the callbacks for checkpoints and early stopping
```

```
cp = tf.keras.callbacks.ModelCheckpoint(filepath="autoencoder_fraud.h5",
                                         mode='min', monitor='val_loss', verbose=2, save_best_only=True)
```

```
# define our early stopping
```

```
early_stop = tf.keras.callbacks.EarlyStopping(
    monitor='val_loss',
    min_delta=0.0001,
    patience=10,
    verbose=1,
    mode='min',
    restore_best_weights=True)
```

```
autoencoder.compile(metrics=['accuracy'],
                     loss='mean_squared_error',
                     optimizer='adam')
```

```
history = autoencoder.fit(normal_train_data, normal_train_data,
                          epochs=nb_epoch,
                          batch_size=batch_size,
                          shuffle=True,
                          validation_data=(test_data, test_data),
                          verbose=1,
                          callbacks=[cp, early_stop]
                          ).history
```

```
Epoch 1/50
```

```
3551/3554 [=====>.] - ETA: 0s - loss: 3.6159e-04 - accuracy: 0.0000
```

```
Epoch 1: val_loss improved from inf to 0.00005, saving model to autoencoder_fraud.h5
```

```
3554/3554 [=====] - 10s 3ms/step - loss: 3.6132e-04 - accuracy: 0.0000
```

```
Epoch 2/50
```

```
3549/3554 [=====>.] - ETA: 0s - loss: 3.8823e-05 - accuracy: 0.0000
```

```
Epoch 2: val_loss did not improve from 0.00005
```

```
3554/3554 [=====] - 9s 3ms/step - loss: 3.8809e-05 - accuracy: 0.0000
```

```
Epoch 3/50
```

```
3552/3554 [=====>.] - ETA: 0s - loss: 3.1846e-05 - accuracy: 0.0000
```

```
Epoch 3: val_loss did not improve from 0.00005
```

```
3554/3554 [=====] - 9s 3ms/step - loss: 3.1843e-05 - accuracy: 0.0000
```

```
Epoch 4/50
```

```
3535/3554 [=====>.] - ETA: 0s - loss: 2.9633e-05 - accuracy: 0.0000
```

```
Epoch 4: val_loss did not improve from 0.00005
```

```
3554/3554 [=====] - 9s 3ms/step - loss: 2.9654e-05 - accuracy: 0.0000
```

```
Epoch 5/50
```

```
3543/3554 [=====>.] - ETA: 0s - loss: 2.5502e-05 - accuracy: 0.0000
```

```

Epoch 5: val_loss did not improve from 0.00005
3554/3554 [=====] - 9s 3ms/step - loss: 2.5484e-05 - accuracy: 0.9999
Epoch 6/50
3547/3554 [=====>.] - ETA: 0s - loss: 2.5053e-05 - accuracy: 0.9999
Epoch 6: val_loss did not improve from 0.00005
3554/3554 [=====] - 9s 3ms/step - loss: 2.5034e-05 - accuracy: 0.9999
Epoch 7/50
3532/3554 [=====>.] - ETA: 0s - loss: 2.3355e-05 - accuracy: 0.9999
Epoch 7: val_loss did not improve from 0.00005
3554/3554 [=====] - 12s 3ms/step - loss: 2.3293e-05 - accuracy: 0.9999
Epoch 8/50
3541/3554 [=====>.] - ETA: 0s - loss: 2.2482e-05 - accuracy: 0.9999
Epoch 8: val_loss did not improve from 0.00005
3554/3554 [=====] - 9s 3ms/step - loss: 2.2503e-05 - accuracy: 0.9999
Epoch 9/50
3551/3554 [=====>.] - ETA: 0s - loss: 2.2096e-05 - accuracy: 0.9999
Epoch 9: val_loss did not improve from 0.00005
3554/3554 [=====] - 12s 3ms/step - loss: 2.2102e-05 - accuracy: 0.9999
Epoch 10/50
3540/3554 [=====>.] - ETA: 0s - loss: 2.1491e-05 - accuracy: 0.9999
Epoch 10: val_loss did not improve from 0.00005
3554/3554 [=====] - 13s 4ms/step - loss: 2.1454e-05 - accuracy: 0.9999
Epoch 11/50
3551/3554 [=====>.] - ETA: 0s - loss: 2.1228e-05 - accuracy: 0.9999
Epoch 11: val_loss did not improve from 0.00005
Restoring model weights from the end of the best epoch: 1.
3554/3554 [=====] - 10s 3ms/step - loss: 2.1229e-05 - accuracy: 0.9999
Epoch 11: early stopping

```

```
# Detect Anomalies on test data
```

```

test_x_predictions = autoencoder.predict(test_data)
mse = np.mean(np.power(test_data - test_x_predictions, 2), axis=1)
error_df = pd.DataFrame({'Reconstruction_error': mse,
                        'True_class': test_labels})


```

```
1781/1781 [=====] - 2s 1ms/step
```

```

threshold_fixed = 52
pred_y = [1 if e > threshold_fixed else 0 for e in error_df.Reconstruction_error.values]
error_df['pred'] = pred_y
conf_matrix = confusion_matrix(error_df.True_class, pred_y)
plt.figure(figsize=(4, 4))
sns.heatmap(conf_matrix, xticklabels=LABELS, yticklabels=LABELS, annot=True, fmt="d");
plt.title("Confusion matrix")
plt.ylabel('True class')
plt.xlabel('Predicted class')
plt.show()

```



[Colab paid products](#) - [Cancel contracts here](#)

